

# **Managing Multiuser Databases**

**Database Processing:  
Fundamentals, Design, and Implementation**  
David M. Kroenke's

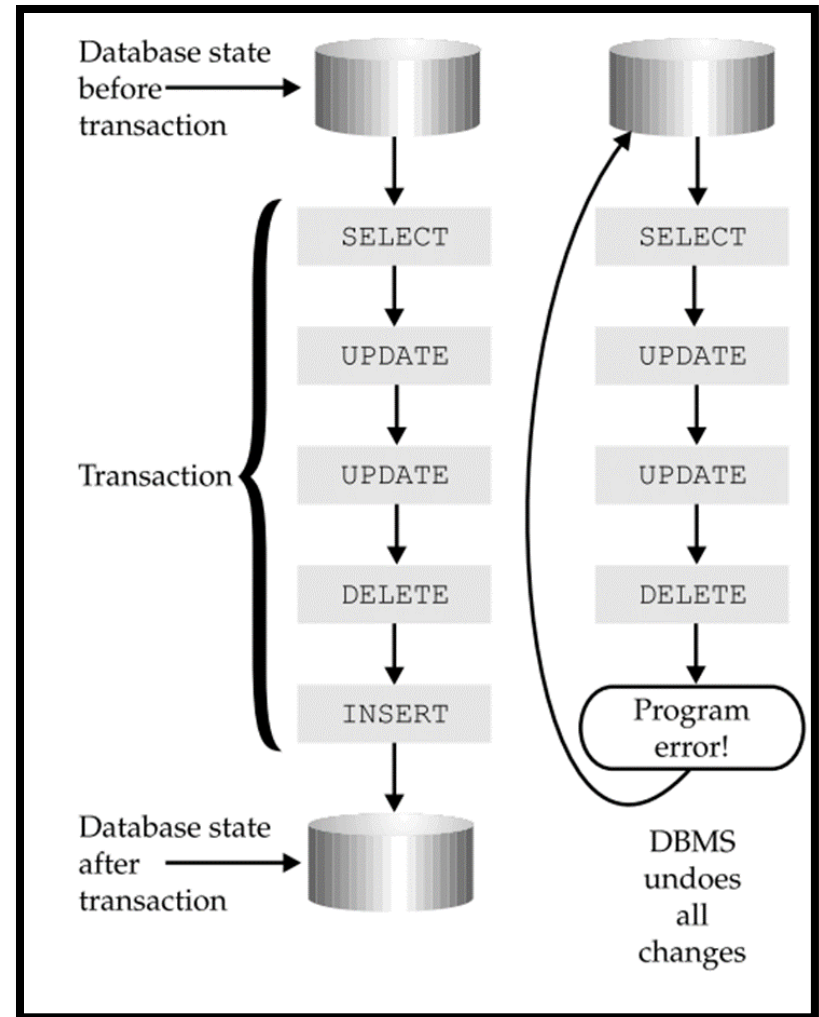
**Oct-Nov, 2023**

# Concurrency Control

- † Most challenging problem in the multiuser DB environment
- † Many users attempt to access the same data resource simultaneously and each user thinks that he/she has sole use of the computer system
- † **Concurrency control** ensures that one user's work does not inappropriately influence another user's work
  - No single concurrency control technique is ideal for all circumstances
  - Trade-offs need to be made between level of protection and throughput

# Transactions

- † Each user submits work in units called transaction
- † Also called “Atomic transactions” or “Logical Unit of Work (LUW)”
- † Each transaction is a series of actions taken against the database that occurs as an atomic unit
  - Either all actions in a transaction occur or none of them do



# Transactions

- † Ex: transaction to enter customer order actions:
  1. Change customer record with new order data
  2. Change salesperson record with new order data
  3. Insert a new order record into the database
  
- † Now visualize the errors introduced without atomic transaction

# Errors Introduced Without Atomic Transaction

Before

Action

After

## CUSTOMER

CNum	OrderNum	Description	AmtDue
123	1000	400 Baseballs	\$2400

## SALESPERSON

Name	Total-Sales	Commission Due
JONES	\$3200	\$320

## ORDER

OrderNum	
1000	...
2000	...
3000	...
4000	...
5000	...
6000	...
7000	...

\*FULL\*

- START
1. Add new-order data to CUSTOMER.
  2. Add new-order data to SALESPERSON.
  3. Insert new ORDER.
- STOP

## CUSTOMER

CNum	OrderNum	Description	AmtDue
123	1000	400 Baseballs	\$2400
123	8000	250 Basketballs	\$6500

## SALESPERSON

Name	Total-Sales	Commission Due
JONES	\$9700	\$970

## ORDER

OrderNum	
1000	...
2000	...
3000	...
4000	...
5000	...
6000	...
7000	...

\*FULL\*

Courtesy: David M. Kronke



# Errors Prevented With Atomic Transaction

## Before

### CUSTOMER

CNum	OrderNum	Description	AmtDue
123	1000	400 Baseballs	\$2400

### SALESPERSON

Name	Total-Sales	Commission Due
JONES	\$3200	\$320

### ORDER

OrderNum	
1000	...
2000	...
3000	...
4000	...
5000	...
6000	...
7000	...

\*FULL\*

## Transaction

Begin Transaction  
Change CUSTOMER data  
Change SALESPERSON data  
Insert ORDER data  
If no errors then  
Commit Transactions  
Else  
Rollback Transaction  
End If

The commands  
marking  
transaction logic  
boundaries  
must be issued  
by the  
application  
program

## After

### CUSTOMER

CNum	OrderNum	Description	AmtDue
123	1000	400 Baseballs	\$2400

### SALESPERSON

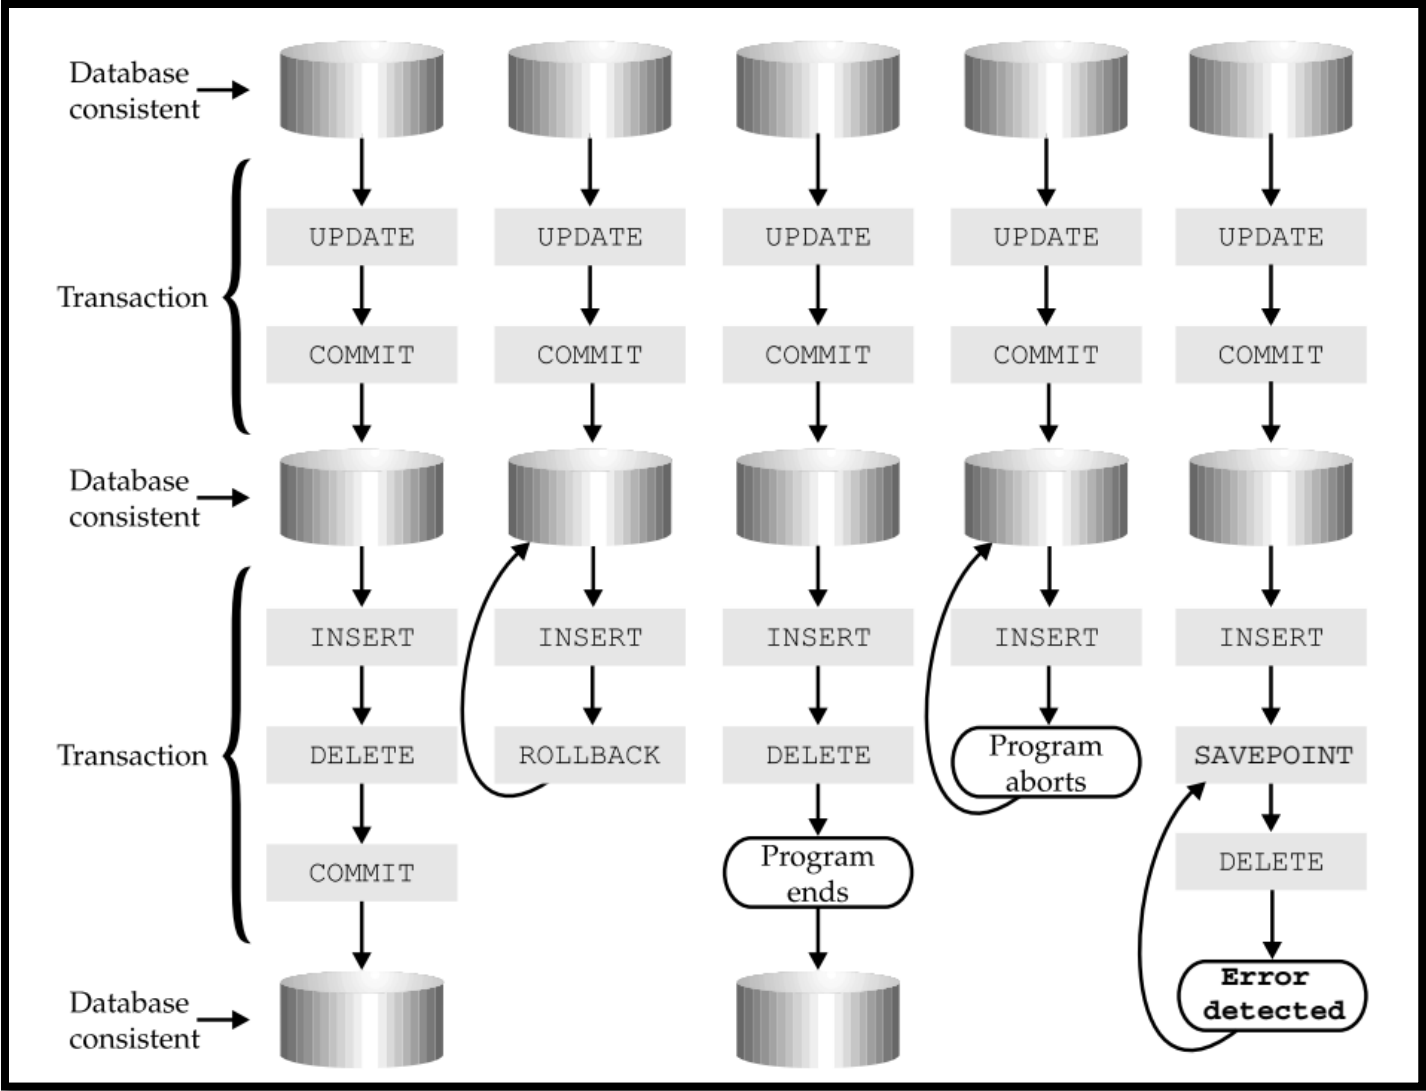
Name	Total-Sales	Commission Due
JONES	\$3200	\$320

### ORDER

OrderNum	
1000	...
2000	...
3000	...
4000	...
5000	...
6000	...
7000	...

\*FULL\*

## Committed and Rolled Back Transactions



# Concurrent Transaction

- † **Concurrent transactions** refer to two or more transactions that appear to users as they are being processed against a database at the same time
- † In reality, CPU can execute only 1 instruction at a time
  - **Transactions are interleaved** meaning that the operating system quickly switches CPU services among tasks so that some portion of each of them is carried out in a given interval
- † Concurrency problems:
  - lost update or inconsistent reads



# Concurrent Transaction Processing

## User A

1. Read item 100
2. Change item 100
3. Write item 100

## User B

1. Read item 200
2. Change item 200
3. Write item 200

## Order of Processing at DB server

1. Read item 100 for A
2. Read item 200 for B
3. Change item 100 for A
4. Write item 100 for A
5. Change item 200 for B
6. Write item 200 for B

# Lost-Update Problem

## User A

1. Read item 100 (Item count is 10)
2. Reduce count of items by 5
3. Write item 100

## User B

1. Read item 100 (Item count is 10)
2. Reduce count of items by 3
3. Write item 100

## Order of Processing at DB server

1. Read item 100 from A (item count = 10)
2. Read item 100 from B (item count = 10)
3. Set item count to 5 (for A)
4. Write item 100 for A
5. Set item count to 7 (for B)
6. Write item 100 for B

concurrent  
update  
problem

\* Note the change and write in step 3 and 4 are lost

# Inconsistent Read Problem

- † Both users obtained current data from the DB
- † User B reads data that have been processes by a portion of a transaction from User A
- † As a result User B reads incorrect data
- † Solution: prevent multiple applications from having copies of the same record when the record is about to be changed
- † Resource Locking

# Resource Locking

- † **Resource locking** prevents multiple applications from obtaining copies of the same record when the record is about to be changed
- † Disallow sharing by locking data that are retrieved for update
- † Use **Locks**

# Concurrent Processing with Explicit Locks

## User A

1. Lock item 100
2. Read item 100
3. Reduce count by 5
4. Write item 100

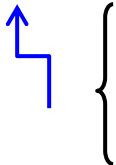
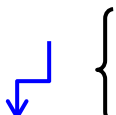
## User B

1. Lock item 100
2. Read item 100
3. Reduce count by 3
4. Write item 100

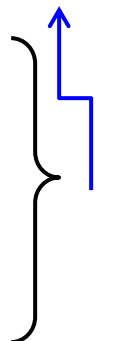
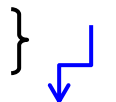
## Order of Processing at DB server

1. Lock item 100 for A
2. Read item 100 from A
3. Lock item 100 for B; cannot so place B in wait state
4. set item count to 5 (for A)
5. Write item 100 for A
6. Release A's lock on item 100
7. Place lock on item 100 for B
8. Read item 100 for B
9. set item count to 2 for B
10. Write item 100 for B
11. Release B's lock on item 100

A's  
Transaction



B's  
Transaction





# Lock Terminology

- † **Implicit locks** are locks placed by the DBMS
- † **Explicit locks** are issued by the application program
- † **Lock granularity** refers to size of a locked resource: rows, page, table, and database level
- † Large granularity:
  - easy to manage but frequently causes conflicts
- † Types of locks
  - An **exclusive lock** locks the item from access of any type. No other transaction can read or change the data
  - A **shared lock** allows other users to read the locked resource, but they cannot update it

# Serializable Transactions

- † **Serializable Transactions** refer to two transactions that run concurrently and generate results that are consistent with the results that would have occurred if they had run separately
- † **Two-phased locking** is one of the techniques used to achieve serializability
  - **Growing phase:** Transactions are allowed to obtain locks as necessary
  - **Shrinking phase:** The locks held by the transaction are being released
  - Once the 1<sup>st</sup> lock is released, no other lock can be obtained

# Strict Two-phased Locking

- † A special case of two-phased locking
  - Locks are obtained throughout the transaction
  - No lock is released until the COMMIT or ROLLBACK command is issued
  - This strategy is more restrictive but easier to implement than two-phased locking

# Deadlock

† Locking solves one problem but introduces another

## User A

1. Lock Paper
2. Take Paper
3. Lock Pencils

Deadly  
Embrace

## User B

1. Lock Pencils
2. Take Pencils
3. Lock Paper

## Order of Processing at DB server

1. Lock Paper for User A
  2. Lock Pencils for User B
  3. Process A's request; write paper record
  4. Process B's request; write pencil record
  5. Put A in wait state for Pencils
  6. Put B in wait state for Paper
- \*\*\*\*Locked\*\*\*\*



# Deadlock

- † **Deadlock** occurs when two transactions are each waiting on a resource that the other transaction holds
- † 3 common ways to solve this situation
  - Deadlock Prevention
    - Allow users to issue all lock requests at one time
  - Deadlock Avoidance
    - Require all application programs to lock resources in some predetermined order
  - Breaking deadlock
    - Almost every DBMS has algos for detecting deadlock
    - When deadlock occurs, DBMS aborts one of the transactions and rollbacks partially completed work



# Dimensions of Resource Locking

- † Lock level (Lock granularity)
- † Lock scope
  - Number of granules to be locked
- † Lock duration
  - To end of job or terminal session (maximum)
  - For single DBMS action (minimum)
  - To end of transaction
- † Lock agent (responsibility)
  - DBMS function (more reliable, less flexible)
  - Application program function (less reliable, more flexible)

# Optimistic versus Pessimistic Locking (two basic styles)

- † **Optimistic locking** assumes that no transaction conflict will occur:
  - DBMS processes a transaction, updates are issued, and then checks whether conflict occurred:
    - If not, the transaction is finished
    - If so, the transaction is repeated until there is no conflict
- † **Pessimistic locking** assumes that conflict will occur:
  - Locks are issued before a transaction is processed, and then the locks are released

# Optimistic Locking

SELECT PRODUCT.Name, PRODUCT.Quantity  
FROM PRODUCT  
WHERE PRODUCT.Name = 'Pencil'

Set NewQuantity = PRODUCT.Quantity - 5

{process transaction – take exception action if NewQuantity < 0, etc.

Assuming all is OK: }

LOCK PRODUCT

**Granularity has to be decided**

UPDATE PRODUCT  
SET PRODUCT.Quantity = NewQuantity  
WHERE PRODUCT.Name = 'Pencil'  
AND PRODUCT.Quantity = OldQuantity

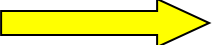
UNLOCK PRODUCT

{check to see if update was successful;  
if not, repeat transaction}

# Pessimistic Locking

**Granularity has to be decided**

 LOCK      PRODUCT

 {  
SELECT    PRODUCT.Name, PRODUCT.Quantity  
FROM      PRODUCT  
WHERE     PRODUCT.Name = 'Pencil'

Set NewQuantity = PRODUCT.Quantity - 5

{process transaction - take exception action if NewQuantity < 0, etc.

Assuming all is OK: }

UPDATE    PRODUCT  
SET        PRODUCT.Quantity = NewQuantity  
WHERE     PRODUCT.Name = 'Pencil'

 UNLOCK    PRODUCT

{no need to check if update was successful}

# Optimistic Locking (+ve and -ve points)

- † The lock is held for less time than with pessimistic locking
  - More imp in case lock granularity is large
- † Transaction may have to be repeated many times
  - Transactions that involve a lot of activity on a given row are poorly suited
- † Optimistic locking is preferred for the Internet and for many intranet applications



# Declaring Lock Characteristics

- † Most application programs do not explicitly declare locks due to its complication
- † They mark transaction boundaries and declare locking behavior they want the DBMS to use
  - Transaction boundary markers: **BEGIN, COMMIT, and ROLLBACK TRANSACTION**
- † Advantage:
  - If the locking behavior needs to be changed, only the lock declaration need be changed, not the application program

# Marking Transaction Boundaries

 BEGIN TRANSACTION:

```
SELECT    PRODUCT.Name, PRODUCT.Quantity
FROM      PRODUCT
WHERE     PRODUCT.Name = 'Pencil'
```

Old Quantity = PRODUCT.Quantity

Set NewQuantity = PRODUCT.Quantity - 5

{process transaction – take exception action if NewQuantity < 0, etc.}

```
UPDATE    PRODUCT
SET       PRODUCT.Quantity = NewQuantity
WHERE     PRODUCT.Name = 'Pencil'
```

{continue processing transaction} . . .

IF transaction has completed normally      THEN

 COMMIT TRANSACTION

ELSE

 ROLLBACK TRANSACTION

END IF

Continue processing other actions not part of this transaction . . .

† These boundaries are the essential information that the DBMS needs in order to enforce different locking strategies

# Marking Transaction Boundaries

BEGIN TRANSACTION:

```
SELECT    PRODUCT.Name, PRODUCT.Quantity
FROM      PRODUCT
WHERE     PRODUCT.Name = 'Pencil'
```

Old Quantity = PRODUCT.Quantity

Set NewQuantity = PRODUCT.Quantity - 5

{process transaction – take exception action if NewQuantity < 0, etc.}

```
UPDATE    PRODUCT
SET       PRODUCT.Quantity = NewQuantity
WHERE     PRODUCT.Name = 'Pencil'
```

{continue processing transaction} . . .

IF transaction has completed normally      THEN

    COMMIT TRANSACTION

ELSE

    ROLLBACK TRANSACTION

END IF

Continue processing other actions not part of this transaction . . .

† If developer now declares (via a system parameter or similar means) that he/she wants optimistic locking, the DBMS will implicitly set the locks in the correct place for that locking style

# ACID Transactions

† An **ACID** transaction is one that is:

**A**tomic,

**C**onsistent,

**I**solated, and

**D**urable

† **Atomic** means either all of the database actions occur or none of them do

† **Durable** means all committed changes are permanent



# ACID Transactions (Consistency)

- † **Consistency** means that each user sees a consistent view of the data, including visible changes made by the user's own transactions and transactions of other users
- † **Consistency:**
  - **Statement level consistency**
  - **Transaction level consistency**



# ACID Transactions (Consistency)

## † **Statement level consistency:**

- Each statement independently processes rows consistently; but
- Changes from other users to these rows might be allowed during the interval between the two SQL statements

## † **Transaction level consistency:**

- All rows impacted by either of the SQL statements are protected from changes during the entire transaction
  - With transaction level consistency, a transaction may not see its own changes

# ACID Transactions (Consistency)

† Statement level consistency

† Transaction level consistency

- Second SQL statement may not see rows changed by the first SQL statement

```
UPDATE CUSTOMER
SET      Areacode = '425'
WHERE    Zipcode = '98050'
COMMIT TRANSACTION
```

```
BEGIN TRANSACTION
UPDATE CUSTOMER
SET      Areacode = '425'
WHERE    Zipcode = '98050'

{OTHER TRANSACTION WORK}

UPDATE CUSTOMER
SET      Discount = 0.05
WHERE    Areacode = '425'

{OTHER TRANSACTION WORK}

COMMIT TRANSACTION
```

# Transactions (isolation)

- † DBMS must ensure that the user's actions do not interfere with one another
- † Four fundamental problems to be addressed:
  - **Lost update**
  - **Dirty read**
  - **Nonrepeatable reads**
  - **Phantom reads**

# ACID Transactions (isolation)

## † Dirty read: (Uncommitted Data Problem)

- One transaction reads a changed record that has not been committed to the db
  - If one transaction reads a row changed by a second uncommitted transaction, and this second transaction later aborts
  - The danger of a dirty read is that the uncommitted change can be rolled back. If so, the transaction that made the dirty read will process incorrect data.

Joe's Program

PRODUCTS Table

Mary's Program

12:00

MFR_ID	PRODUCT_ID	QTY_ON_HAND
•		
•		
•		
ACI	41004	139
•		
•		
•		

12:01

```
SELECT QTY_ON_HAND
FROM PRODUCTS...
```

Answer: 139

Accept order for 100

12:04

```
UPDATE PRODUCTS
SET QTY_ON_HAND = 39...
```

12:06

ROLLBACK

12:04

MFR_ID	PRODUCT_ID	QTY_ON_HAND
•		
•		
•		
ACI	41004	39
•		
•		
•		

12:06

MFR_ID	PRODUCT_ID	QTY_ON_HAND
•		
•		
•		
ACI	41004	139
•		
•		
•		

12:05

```
SELECT QTY_ON_HAND
FROM PRODUCTS
```

Answer: 39

Refuse order for 125

Note: Buy more!!!

**Uncommitted  
Data Problem**



# ACID Transactions (isolation)

- † Nonrepeatable reads: (Inconsistent Data Problem)
  - When a transaction rereads data it has previously read and finds modification or deletions caused by a committed transaction

12:00

MFR_ID	PRODUCT_ID	QTY_ON_HAND
•		
•		
•		
ACI	41004	139
•		
•		
•		

12:01

```
SELECT QTY_ON_HAND
FROM PRODUCTS...
```

Answer: 139

Accept order for 100

12:04

```
UPDATE PRODUCTS
SET QTY_ON_HAND = 39...
```

12:04

MFR_ID	PRODUCT_ID	QTY_ON_HAND
•		
•		
•		
ACI	41004	39
•		
•		
•		

12:02

```
SELECT QTY_ON_HAND
FROM PRODUCTS...
...ID = 41004
```

Answer: 139

12:03

```
SELECT QTY_ON_HAND
FROM PRODUCTS...
...ID = 41005
```

12:05

```
SELECT QTY_ON_HAND
FROM PRODUCTS
...ID = 41004
```

Answer: 39

**Inconsistent  
Data Problem**

# ACID Transactions (isolation)

† Phantom insert:

- When a transaction rereads data and finds new rows that were inserted by a committed transaction since the prior read

## Update Program

## ORDERS Table

12:00

ORDER_NUM	AMOUNT
112961	\$31,500.00
113012	\$3,745.00
•	
•	
•	

12:04

```
INSERT INTO ORDERS VALUES  
(118102, .....5,000.00)
```

12:05

```
COMMIT
```

12:04

ORDER_NUM	AMOUNT
112961	\$31,500.00
118102	\$5,000.00
113012	\$3,745.00
•	
•	
•	

## Report Program

12:00

```
SELECT *  
FROM ORDERS
```

12:01

Answer:  
112961, \$31,500

12:02

Answer:  
113012, \$3,745

•  
•  
•

12:10

```
SELECT *  
FROM ORDERS
```

12:11

Answer:  
112961, \$31,500

12:12

Answer:  
118102, \$5,000

12:13

Answer:  
113012, \$3,745

•  
•

**Phantom  
Insert Problem**

# ACID Transactions (isolation)

- † Under a strict definition of a SQL transaction, no action by a concurrently executing transaction is allowed to impact the data visible during the course of your transaction
- † If your program performs a DB query during a transaction, proceeds with other work, and later performs the same DB query a second time, the SQL transaction mechanism guarantees that the data returned by the two queries will be identical (unless your transaction acted to change the data)



# ACID Transactions (isolation)

- † Absolute isolation of your transaction from all other concurrently executing transaction is very costly in terms of DB locking and loss of DB concurrency
- † **Isolation** means application programmers are able to declare the type of isolation level and to have the DBMS manage locks so as to achieve that level of isolation
- † SQL 2 defines four transaction isolation levels:
  - **Read uncommitted**
  - **Read committed**
  - **Repeatable read**
  - **Serializable** (default isolation level in ANSI/ISO SQL standard)

# Transaction Isolation Level

Isolation Level	Lost Update	Dirty Read (uncommitted data)	Nonrepeatable Read (inconsistent data)	Phantom Insert
<b>Read Uncommitted</b>	Prevented by DBMS	Can occur	Can occur	Can occur
<b>Read Committed</b>	Prevented by DBMS	Prevented by DBMS	Can occur	Can occur
<b>Repeatable Read</b>	Prevented by DBMS	Prevented by DBMS	Prevented by DBMS	Can occur
<b>Serializable</b>	Prevented by DBMS	Prevented by DBMS	Prevented by DBMS	Prevented by DBMS

† SQL 2 standard defines 4 isolation levels

† Application programmer declares the type of isolation level he/she wants, and DBMS manages locks so as to achieve that level of isolation

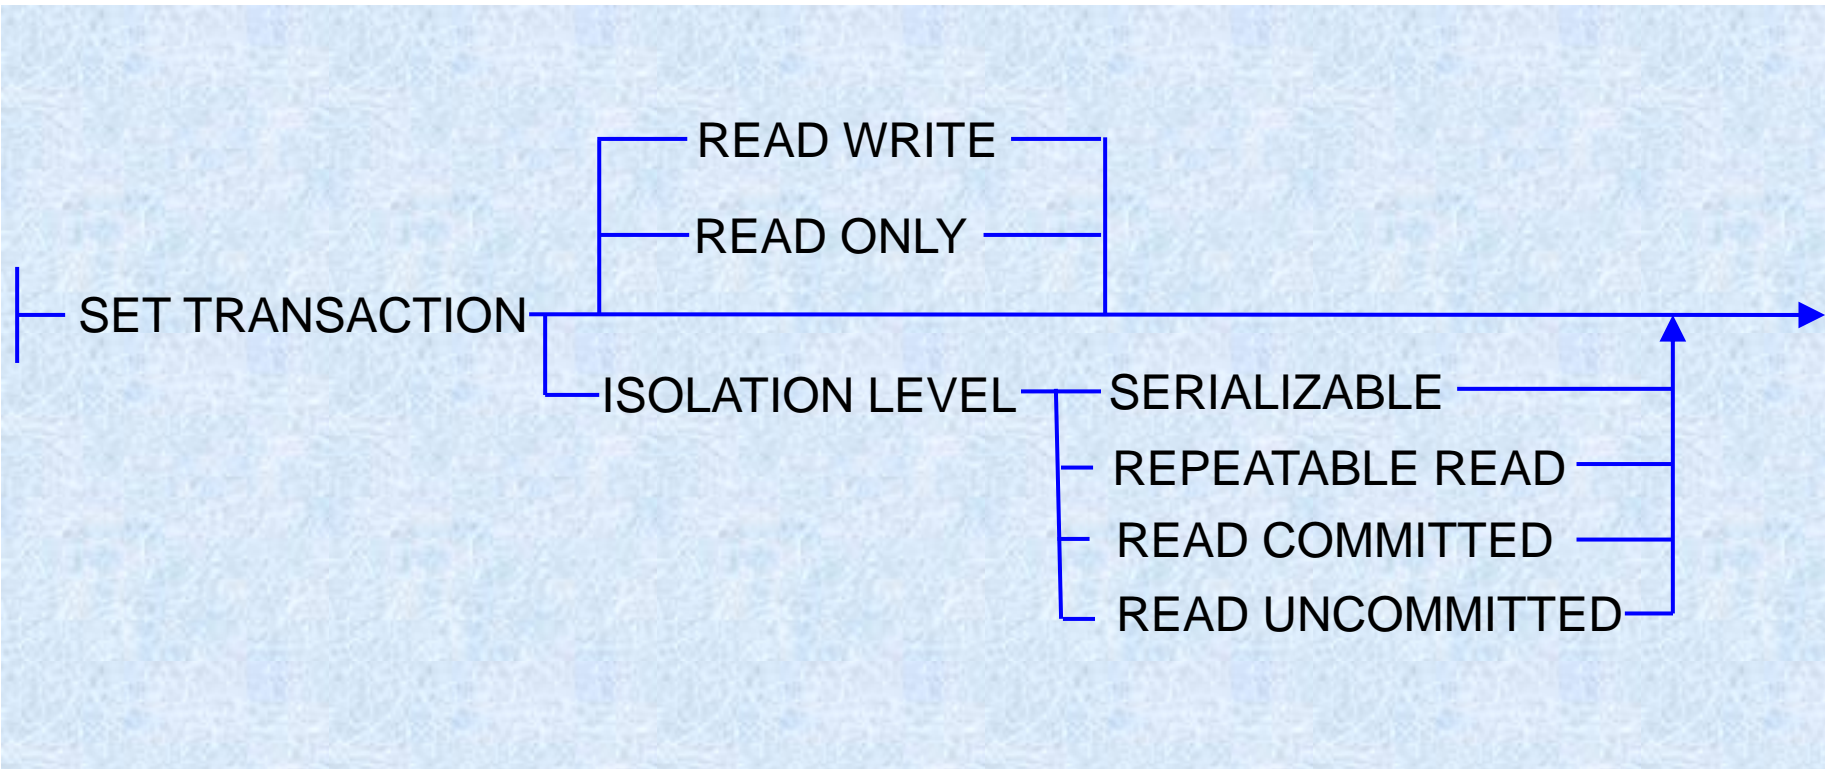
# Transaction Isolation Level

Isolation Level	Lost Update	Dirty Read	Nonrepeatable Read	Phantom Insert
<b>Read Uncommitted</b>	Prevented by DBMS	Can occur	Can occur	Can occur
<b>Read Committed</b>	Prevented by DBMS	Prevented by DBMS	Can occur	Can occur
<b>Repeatable Read</b>	Prevented by DBMS	Prevented by DBMS	Prevented by DBMS	Can occur
<b>Serializable</b>	Prevented by DBMS	Prevented by DBMS	Prevented by DBMS	Prevented by DBMS

† Generally, more restrictive the level, the less the throughput

† Not all DBMS products support all of these levels

# Set Transaction Statement



† <http://www.timlin.net/csm/cis363/tranproc.html> (transaction processing with MySQL)